

## DATABASES

### Chapter 6 : Transaction Management and Concurrency Control

Pr. Kamel BOUKHALFA

Slides From the Textbook :  
Carlos Coronel and Steven Morris, Database Systems: Design, Implementation, and Management  
Tenth Edition

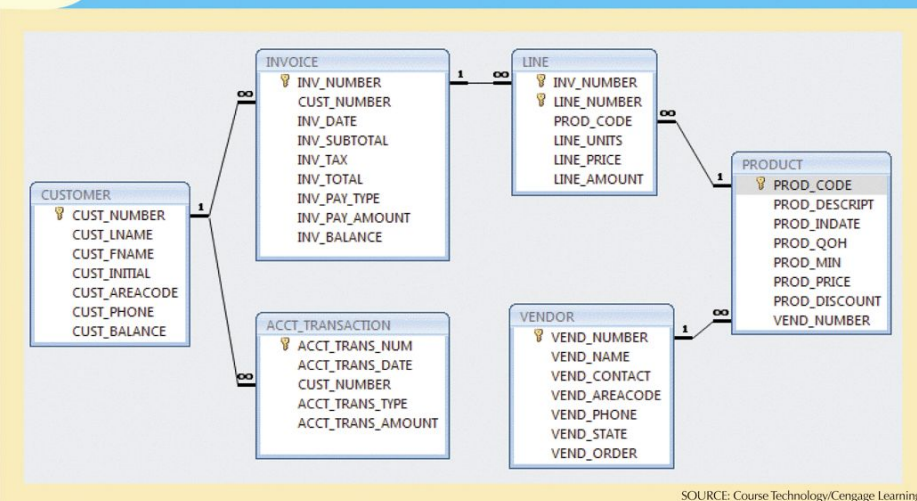
16/10/2022

16/10/2022

2

- ## What Is a Transaction?
- Logical unit of work that must be either entirely completed or aborted
  - Successful transaction changes database from one **consistent state** to another
    - One in which all data integrity constraints are satisfied
  - Most real-world database transactions are formed by two or more database requests
    - Equivalent of a single SQL statement in an application program or transaction
    - Will involve at least one read and write action

FIGURE 10.1 The Ch10\_SaleCo database relational diagram



## Evaluating Transaction Results

- Not all transactions update database
- SQL code represents a transaction because database was accessed
- Improper or incomplete transactions can have devastating effect on database integrity
  - Some DBMSs provide means by which user can define enforceable constraints
  - Other integrity rules are enforced automatically by the DBMS

# Evaluating Transaction Results

```
INSERT INTO INVOICE VALUES (...);
```

```
INSERT INTO LINE (...);
```

```
UPDATE PRODUCT SET ...
```

```
UPDATE CUSTOMER SET CUST_BALANCE..
```

```
INSERT INTO ACCT_TRANSACTION VALUES (...);
```

```
COMMIT;
```

- Suppose the first 3 statements are executed and then power is lost with no backup –DBMS is in an inconsistent state unless transaction management is supported
  - Invoice and Line rows were added and the QOH of Product table updated but the customer balance has not been increased nor has a new transaction been added to Acct\_Transaction

FIGURE 10.2 Tracing the transaction in the Ch10\_SaleCo database

Database name: Ch10\_SaleCo

Table name: INVOICE

| INV_NUMBER | CUST_NUMBER | INV_DATE  | INV_SUBTOTAL | INV_TAX | INV_TOTAL | INV_PAY_TYPE | INV_PAY_AMOUNT | INV_BALANCE |
|------------|-------------|-----------|--------------|---------|-----------|--------------|----------------|-------------|
| 1001       | 10014       | 18-Jan-12 | 54.92        | 4.39    | 59.31     | cc           | 59.31          | 0.00        |
| 1002       | 10011       | 18-Jan-12 | 9.99         | 0.80    | 10.79     | cash         | 10.79          | 0.00        |
| 1003       | 10012       | 18-Jan-12 | 270.70       | 21.66   | 292.36    | cc           | 292.36         | 0.00        |
| 1004       | 10011       | 17-Jan-12 | 34.87        | 2.79    | 37.66     | cc           | 37.66          | 0.00        |
| 1005       | 10016       | 17-Jan-12 | 79.44        | 5.44    | 84.88     | cc           | 84.88          | 0.00        |
| 1006       | 10014       | 17-Jan-12 | 397.83       | 31.83   | 429.66    | credit       | 100.00         | 329.66      |
| 1007       | 10016       | 17-Jan-12 | 34.97        | 2.80    | 37.77     | chq          | 37.77          | 0.00        |
| 1008       | 10011       | 17-Jan-12 | 1033.00      | 82.65   | 1115.72   | credit       | 500.00         | 615.72      |
| 1009       | 10016       | 18-Jan-12 | 256.99       | 20.56   | 277.55    | credit       | 0.00           | 277.55      |

Table name: LINE

| INV_NUMBER | LINE_NUMBER | PROD_CODE | LINE_UNITS | LINE_PRICE | LINE_AMOUNT |
|------------|-------------|-----------|------------|------------|-------------|
| 1001       | 1           | 13-G2P2   | 3          | 14.99      | 44.97       |
| 1001       | 2           | 23109-HB  | 1          | 9.95       | 9.95        |
| 1002       | 1           | 54778-2T  | 2          | 4.99       | 9.98        |
| 1003       | 1           | 22369-GPD | 4          | 38.95      | 155.80      |
| 1003       | 2           | 1546-GC2  | 1          | 39.95      | 39.95       |
| 1003       | 3           | 13-G2P2   | 5          | 14.99      | 74.95       |
| 1004       | 1           | 54778-2T  | 3          | 4.99       | 14.97       |
| 1004       | 2           | 23109-HB  | 2          | 9.95       | 19.90       |
| 1005       | 1           | PVC3104RT | 12         | 5.67       | 70.44       |
| 1006       | 1           | 394-107T  | 3          | 6.99       | 20.97       |
| 1006       | 2           | 2332Q-T1  | 1          | 109.92     | 109.92      |
| 1006       | 3           | 23109-HB  | 1          | 9.95       | 9.95        |
| 1006       | 4           | 89-VWSE-G | 1          | 256.99     | 256.99      |
| 1007       | 1           | 13-G2P2   | 2          | 14.99      | 29.98       |
| 1007       | 2           | 54778-2T  | 1          | 4.99       | 4.99        |
| 1007       | 1           | PVC3104RT | 5          | 5.67       | 28.35       |
| 1008       | 2           | WR37T3    | 4          | 119.95     | 479.80      |
| 1009       | 3           | 23109-HB  | 1          | 9.95       | 9.95        |
| 1008       | 4           | 89-VWSE-G | 2          | 256.99     | 513.98      |
| 1009       | 1           | 89-VWSE-G | 1          | 256.99     | 256.99      |

Table name: PRODUCT

| PROD_CODE | PROD_DESCRIPTION                    | PROD_ADJDATE | PROD_QOH | PROD_MN | PROD_PRICE | PROD_DISCOUNT | VEND_NUMBER |
|-----------|-------------------------------------|--------------|----------|---------|------------|---------------|-------------|
| 1156P31   | Power planer, 15 psi, 3-nozzle      | 03-Nov-11    | 8        | 5       | 109.99     | 0.00          | 25595       |
| 13-G2P2   | 7.25-in. pwr. saw blade             | 13-Dec-11    | 32       | 15      | 14.99      | 0.05          | 21344       |
| 14-G1L3   | 9.00-in. pwr. saw blade             | 13-Nov-11    | 18       | 12      | 17.49      | 0.00          | 21344       |
| 1546-GC2  | Hsk. cloth, 14-in., 2-50            | 15-Jan-12    | 15       | 8       | 39.95      | 0.00          | 23119       |
| 1558-0W1  | Hsk. cloth, 12-in., 3-50            | 15-Jan-12    | 23       | 5       | 43.99      | 0.00          | 23119       |
| 2232Q-T1  | B&D gisaw, 12-in. blade             | 30-Dec-11    | 8        | 5       | 109.92     | 0.05          | 24289       |
| 2233QW1E  | B&D gisaw, 8-in. blade              | 24-Dec-11    | 6        | 5       | 99.97      | 0.05          | 24289       |
| 2238GPD   | B&D cordless drill, 1/2-in.         | 20-Jan-12    | 12       | 5       | 38.95      | 0.05          | 25595       |
| 23109-HB  | Claw hammer                         | 20-Jan-12    | 23       | 10      | 9.95       | 0.10          | 21225       |
| 23114-AA  | Sledge hammer, 12 lb.               | 02-Jan-12    | 8        | 5       | 14.40      | 0.05          |             |
| 54778-2T  | Rat-tail file, 1.6-in. file         | 15-Dec-11    | 43       | 20      | 4.99       | 0.00          | 21344       |
| 89-VWSE-G | Hood chain saw, 18 in.              | 07-Jan-12    | 11       | 5       | 256.99     | 0.05          | 24289       |
| PVC3104RT | PVC pipe, 3.5-in., 8-ft             | 06-Jan-12    | 188      | 75      | 5.67       | 0.00          |             |
| SM-18277  | 1.25-in. metal screw, 25            | 01-Mar-12    | 172      | 75      | 6.99       | 0.00          | 21225       |
| SW-23116  | 2.5-in. wtd. screw, 50              | 24-Feb-12    | 237      | 100     | 8.45       | 0.00          | 21231       |
| W837T3    | Steel nutting, 4/16 x 1/8", 5" mesh | 17-Jan-12    | 16       | 5       | 119.95     | 0.10          | 25595       |

Table name: CUSTOMER

| CUST_NUM | CUST_LNAME | CUST_FNAME | CUST_INITIAL | CUST_AREACODE | CUST_PHONE | CUST_BALANCE |
|----------|------------|------------|--------------|---------------|------------|--------------|
| 10010    | Ramas      | Alfred     | A            | 615           | 844-2573   | 0.00         |
| 10011    | Dunne      | Leena      | K            | 713           | 894-1238   | 615.73       |
| 10012    | Smith      | Kathy      | W            | 615           | 894-2265   | 0.00         |
| 10013    | Olowski    | Paul       | F            | 615           | 894-2180   | 0.00         |
| 10014    | Orlando    | Myron      |              | 615           | 222-1672   | 0.00         |
| 10015    | O'Brien    | Adm        | B            | 713           | 442-3361   | 0.00         |
| 10016    | Brown      | James      | G            | 615           | 267-1228   | 277.55       |
| 10017    | Williams   | George     |              | 615           | 260-2556   | 0.00         |
| 10018    | Farnes     | Anne       | G            | 713           | 362-7185   | 0.00         |
| 10019    | Smith      | Owne       | J            | 615           | 267-3809   | 0.00         |

Table name: ACCT\_TRANSACTION

| ACCT_TRANS_NUM | ACCT_TRANS_DATE | CUST_NUMBER | ACCT_TRANS_TYPE | ACCT_TRANS_AMOUNT |
|----------------|-----------------|-------------|-----------------|-------------------|
| 10003          | 17-Jan-12       | 10014       | charge          | 329.66            |
| 10004          | 17-Jan-12       | 10011       | charge          | 615.73            |
| 10006          | 28-Jan-12       | 10014       | payment         | 329.66            |
| 10007          | 18-Jan-12       | 10016       | charge          | 277.55            |

Figure 9.2

SOURCE: Course Technology/Cengage Learning

## Transaction Properties

- Atomicity
  - All operations of a transaction must be completed or else the entire transaction is aborted
- Consistency
  - Permanence of database's consistent state
- Isolation
  - Data used during transaction cannot be used by second transaction until the first is completed

## Transaction Properties (cont'd.)

- Durability
  - Once transactions are committed, they cannot be undone
- Serializability
  - Concurrent execution of several transactions yields consistent results
- Multiuser databases are subject to multiple concurrent transactions

# Transaction Management with SQL

TABLE 10.1 A Transaction Log

| TRL_ID | TRX_NUM | PREV_PTR | NEXT_PTR | OPERATION | TABLE                   | ROW_ID   | ATTRIBUTE    | BEFORE_VALUE | AFTER_VALUE |
|--------|---------|----------|----------|-----------|-------------------------|----------|--------------|--------------|-------------|
| 341    | 101     | Null     | 352      | START     | ****Start Transaction   |          |              |              |             |
| 352    | 101     | 341      | 363      | UPDATE    | PRODUCT                 | 1558-QW1 | PROD_QOH     | 25           | 23          |
| 363    | 101     | 352      | 365      | UPDATE    | CUSTOMER                | 10011    | CUST_BALANCE | 525.75       | 615.73      |
| 365    | 101     | 363      | Null     | COMMIT    | **** End of Transaction |          |              |              |             |

↑  
 TRL\_ID = Transaction log record ID  
 TRX\_NUM = Transaction number  
 PTR = Pointer to a transaction log record ID  
 (Note: The transaction number is automatically assigned by the DBMS.)

- Coordination of simultaneous transaction execution in a multiprocessing database
- Objective is to ensure serializability of transactions in a multiuser environment
- The simultaneous execution of transactions over a shared database can create several data integrity and consistency problems
- Three main problems:
  - Lost updates
  - Uncommitted data
  - Inconsistent retrievals

## Lost Updates

- Lost update problem:
  - Two concurrent transactions update same data element

TABLE 10.2 Two Concurrent Transactions to Update QOH

| TRANSACTION            | COMPUTATION                   |
|------------------------|-------------------------------|
| T1: Purchase 100 units | $PROD\_QOH = PROD\_QOH + 100$ |
| T2: Sell 30 units      | $PROD\_QOH = PROD\_QOH - 30$  |

- Table 10.3 shows the serial execution of the transactions under normal circumstances.
- Final result of  $PROD\_QOH = 105$  is correct

TABLE 10.3 Serial Execution of Two Transactions

| TIME | TRANSACTION | STEP                   | STORED VALUE |
|------|-------------|------------------------|--------------|
| 1    | T1          | Read $PROD\_QOH$       | 35           |
| 2    | T1          | $PROD\_QOH = 35 + 100$ |              |
| 3    | T1          | Write $PROD\_QOH$      | 135          |
| 4    | T2          | Read $PROD\_QOH$       | 135          |
| 5    | T2          | $PROD\_QOH = 135 - 30$ |              |
| 6    | T2          | Write $PROD\_QOH$      | 105          |

## Lost Updates

- Table 10.4 shows what happens when a second transaction can read the  $PROD\_QOH$  value of a product before a previous transaction has been committed for that same product
- One of the updates is lost by being overwritten by the other transaction

TABLE 10.4 Lost Updates

| TIME | TRANSACTION | STEP                            | STORED VALUE |
|------|-------------|---------------------------------|--------------|
| 1    | T1          | Read $PROD\_QOH$                | 35           |
| 2    | T2          | Read $PROD\_QOH$                | 35           |
| 3    | T1          | $PROD\_QOH = 35 + 100$          |              |
| 4    | T2          | $PROD\_QOH = 35 - 30$           |              |
| 5    | T1          | Write $PROD\_QOH$ (Lost update) | 135          |
| 6    | T2          | Write $PROD\_QOH$               | 5            |



# Uncommitted Data

- Uncommitted data phenomenon:
  - Two transactions are executed concurrently
  - First transaction rolled back after second already accessed uncommitted data
- Violates the isolation property

**TABLE 10.5** Transactions Creating an Uncommitted Data Problem

| TRANSACTION            | COMPUTATION                             |
|------------------------|---|
| T1: Purchase 100 units | PROD_QOH = PROD_QOH + 100 (Rolled back) |
| T2: Sell 30 units      | PROD_QOH = PROD_QOH - 30                |

**TABLE 10.6** Correct Execution of Two Transactions

| TIME | TRANSACTION | STEP                | STORED VALUE |
|------|-------------|---------------------|--------------|
| 1    | T1          | Read PROD_QOH       | 35           |
| 2    | T1          | PROD_QOH = 35 + 100 |              |
| 3    | T1          | Write PROD_QOH      | 135          |
| 4    | T1          | *****ROLLBACK ***** | 35           |
| 5    | T2          | Read PROD_QOH       | 35           |
| 6    | T2          | PROD_QOH = 35 - 30  |              |
| 7    | T2          | Write PROD_QOH      | 5            |

# Uncommitted Data

- Uncommitted data problem can arise when the ROLLBACK is completed after T2 has begun its execution

**TABLE 10.7** An Uncommitted Data Problem

| TIME | TRANSACTION | STEP                                  | STORED VALUE |
|------|-------------|---------------------------------------|--------------|
| 1    | T1          | Read PROD_QOH                         | 35           |
| 2    | T1          | PROD_QOH = 35 + 100                   |              |
| 3    | T1          | Write PROD_QOH                        | 135          |
| 4    | T2          | Read PROD_QOH (Read uncommitted data) | 135          |
| 5    | T2          | PROD_QOH = 135 - 30                   |              |
| 6    | T1          | ***** ROLLBACK *****                  | 35           |
| 7    | T2          | Write PROD_QOH                        | 105          |

# Inconsistent Retrievals

- Inconsistent retrievals:
  - First transaction accesses data
  - Second transaction alters the data
  - First transaction accesses the data again
- Transaction might read some data before they are changed and other data after changed
- Yields inconsistent results

**TABLE 10.8** Retrieval During Update

| TRANSACTION T1                       | TRANSACTION T2   |
|--------------------------------------|--|
| SELECT SUM(PROD_QOH)<br>FROM PRODUCT | UPDATE PRODUCT<br>SET PROD_QOH = PROD_QOH + 10<br>WHERE PROD_CODE = 1546-QQ2 |
|                                      | UPDATE PRODUCT<br>SET PROD_QOH = PROD_QOH - 10<br>WHERE PROD_CODE = 1558-QW1 |
|                                      | COMMIT;  |

# Inconsistent Retrievals

- A sample of the PRODUCT table, showing the records being updated as specified in the UPDATE code

TABLE 10.9 Transaction Results: Data Entry Correction

|              | BEFORE    | AFTER          |
|--------------|-----------|----------------|
| PROD_CODE    | PROD_QOH  | PROD_QOH       |
| 11QER/31     | 8         | 8              |
| 13-Q2/P2     | 32        | 32             |
| 1546-QQ2     | 15        | (15 + 10) → 25 |
| 1558-QW1     | 23        | (23 - 10) → 13 |
| 2232-QTY     | 8         | 8              |
| 2232-QWE     | 6         | 6              |
| <b>Total</b> | <b>92</b> | <b>92</b>      |

21

# Inconsistent Retrievals

- Inconsistent retrievals are possible during the transaction execution
  - The “after” summation shown below reflects that the value of 25 for product 1546-QQ2 was read after the WRITE statement was completed so the total up to that point is 65
  - The “before” total reflects that the value of 23 for 1558-QQ1 was read before the next WRITE was completed (in which case the value would have been 13 instead of 23)
    - After step 7, the total is 65+23 rather than 65+13
    - The final total will be 102 instead of 92

22

# Inconsistent Retrievals

TABLE 10.10 Inconsistent Retrievals

| TIME | TRANSACTION | ACTION                                    | VALUE | TOTAL       |
|------|-------------|---|-------|-------------|
| 1    | T1          | Read PROD_QOH for PROD_CODE = '11QER/31'  | 8     | 8           |
| 2    | T1          | Read PROD_QOH for PROD_CODE = '13-Q2/P2'  | 32    | 40          |
| 3    | T2          | Read PROD_QOH for PROD_CODE = '1546-QQ2'  | 15    |             |
| 4    | T2          | PROD_QOH = 15 + 10                        |       |             |
| 5    | T2          | Write PROD_QOH for PROD_CODE = '1546-QQ2' | 25    |             |
| 6    | T1          | Read PROD_QOH for PROD_CODE = '1546-QQ2'  | 25    | (After) 65  |
| 7    | T1          | Read PROD_QOH for PROD_CODE = '1558-QW1'  | 23    | (Before) 88 |
| 8    | T2          | Read PROD_QOH for PROD_CODE = '1558-QW1'  | 23    |             |
| 9    | T2          | PROD_QOH = 23 - 10                        |       |             |
| 10   | T2          | Write PROD_QOH for PROD_CODE = '1558-QW1' | 13    |             |
| 11   | T2          | ***** COMMIT *****                        |       |             |
| 12   | T1          | Read PROD_QOH for PROD_CODE = '2232-QTY'  | 8     | 96          |
| 13   | T1          | Read PROD_QOH for PROD_CODE = '2232-QWE'  | 6     | 102         |

23

# The Scheduler

- Special DBMS program
  - Purpose is to establish order of operations within which concurrent transactions are executed
- Interleaves execution of database operations:
  - Ensures serializability
  - Ensures isolation
- Serializable schedule
  - Interleaved execution of transactions yields same results as serial execution
- Transactions that are not serializable with interleaving are executed on a first-come first-serve basis

24

# The Scheduler

- FCFS is an inefficient method of allocating the CPU, since the CPU waits idly as a READ or WRITE operation completes
- The scheduler also facilitates data isolation to ensure that two transactions do not update the same data element at the same time
- If one or both of two concurrent transactions wants to perform a WRITE on the same data, then a conflict will result

TABLE 10.11 Read/Write Conflict Scenarios: Conflicting Database Operations Matrix

| Operations | TRANSACTIONS |       | RESULT      |
|------------|--------------|-------|-------------|
|            | T1           | T2    |             |
| Read       | Read         | Read  | No conflict |
| Read       | Read         | Write | Conflict    |
| Write      | Read         | Read  | Conflict    |
| Write      | Write        | Write | Conflict    |

25

# Concurrency Control with Locking Methods

- Lock
  - Guarantees exclusive use of a data item to a current transaction
    - Lock is acquired prior to access and released when the transaction is completed
  - Required to prevent another transaction from reading inconsistent data
  - Pessimistic locking
    - Use of locks based on the assumption that conflict between transactions is likely
  - Lock manager
    - Responsible for assigning and policing the locks used by transactions

26


## Lock Granularity

- Indicates level of lock use
- Locking can take place at following levels:
  - Database
  - Table
  - Page
  - Row
  - Field (attribute)

## Lock Granularity

- Database-level lock
  - Entire database is locked, prevents the use of any tables by transaction T2 while T1 is being executed
  - Good for batch systems, not for multi-user systems

FIGURE 10.3 Database-level locking sequence



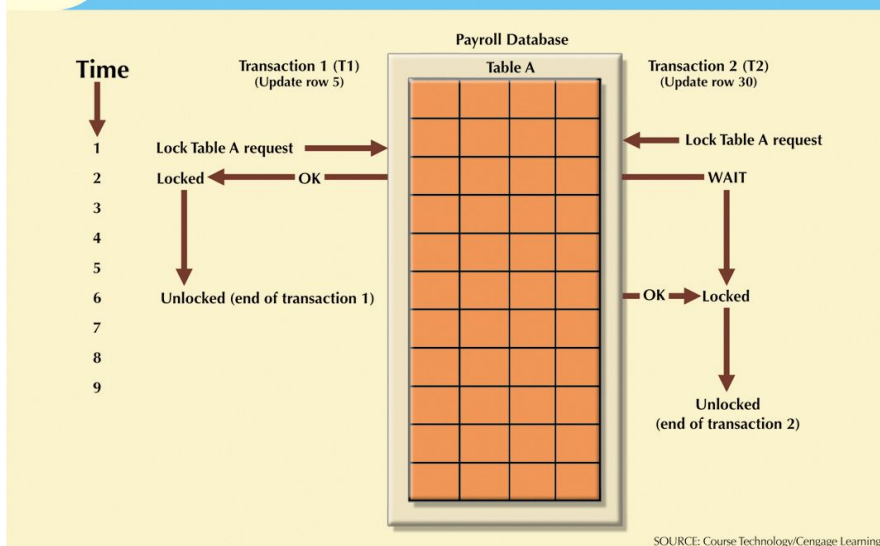
29

Database Systems, 10th Edition

30

- Table-level lock
  - Entire table is locked
  - If a transaction requires access to several tables, each table may be locked
  - Two transactions can access the database as long as they access different tables
  - Not suitable for multi-user systems because all rows are locked even if different transactions access different rows

FIGURE 10.4 An example of a table-level lock

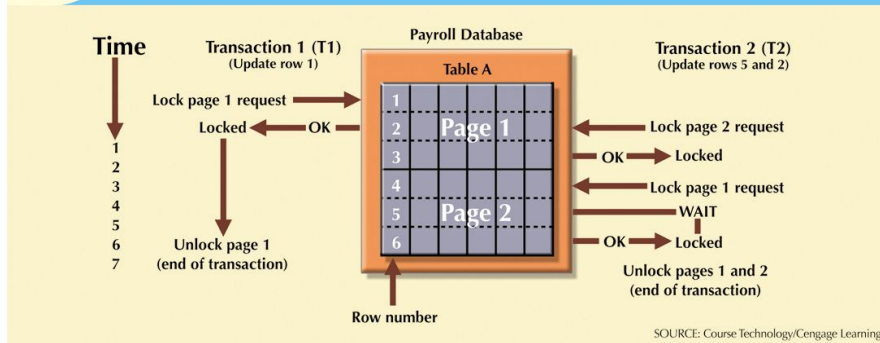


- Page-level lock
  - Entire diskpage (or page) is locked
    - Page is the equivalent of a diskblock - a directly addressable section of disk (e.g., 4K, 8K or 16K)
    - If you want to write 73 bytes to a 4K page, the entire page must be read from disk, updated in memory and written back to disk
      - A table can span multiple pages and a page can contain several rows of one or more tables
  - As long as two transactions do not access the same page, no waiting is required by the second transaction
  - currently the most frequently used locking method for multi-user DBMSs



# Lock Granularity

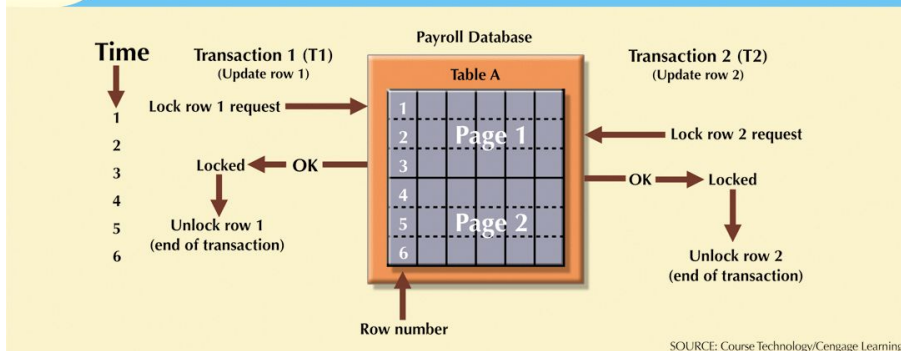
FIGURE 10.5 An example of a page-level lock



- Row-level lock
  - Allows concurrent transactions to access different rows of same table
    - Even if rows are located on same page
    - High overhead because a lock exists for each row of a table
    - If the application requests multiple locks on the same, a row level lock is automatically escalated to a page level lock

# Lock Granularity

FIGURE 10.6 An example of a row-level lock



- Field-level lock
  - Allows concurrent transactions to access same row
    - Requires use of different fields (attributes) within the row
    - Most flexible type of lock but requires an extremely high level of overhead – rarely implemented

## Lock Types

- Binary lock
  - Two states: locked (1) or unlocked (0)
  - Every DB operation requires that the affected object be locked
  - Transaction must unlock the object after its termination
  - Locks and unlocks are performed automatically by the DBMS, not the user

- Applying locks to the “Lost Update Problem”
  - Lock/unlock eliminates problem because the lock is not released until the WRITE is completed
  - PROD\_QOH can not be used until it has been properly updated
  - Binary locks are considered too restrictive as the DBMS will not allow two transactions to read the same object even though neither updates the database and no concurrency problems occur

TABLE 10.12 An Example of a Binary Lock

| TIME | TRANSACTION | STEP                  | STORED VALUE |
|------|-------------|-----------------------|--------------|
| 1    | T1          | Lock PRODUCT          |              |
| 2    | T1          | Read PROD_QOH         | 15           |
| 3    | T1          | $PROD\_QOH = 15 + 10$ |              |
| 4    | T1          | Write PROD_QOH        | 25           |
| 5    | T1          | Unlock PRODUCT        |              |
| 6    | T2          | Lock PRODUCT          |              |
| 7    | T2          | Read PROD_QOH         | 23           |
| 8    | T2          | $PROD\_QOH = 23 - 10$ |              |
| 9    | T2          | Write PROD_QOH        | 13           |
| 10   | T2          | Unlock PRODUCT        |              |

## Lock Types

- Exclusive lock
  - Access is specifically reserved for transaction that locked object
  - Must be used when potential for conflict exists
    - i.e., when two operations are in conflict when they access the same data and at least one of them is performing a WRITE
- Shared lock
  - Concurrent transactions are granted read access on basis of a common lock
    - Produces no conflict as long as all the concurrent transactions are read-only

## Lock Types

- A shared lock is issued when a transaction wants to read data from the DB and no exclusive lock is held on that data item
- An exclusive lock is issued when a transaction wants to update (write) a data item and no locks are currently held on that data item by any other transaction
- Thus, under these concepts, a lock can have three states
  - Unlocked
  - Shared (read)
  - Exclusive (write)

## Lock Types

- If T1 has a shared lock on data item X, T2 may be issued a shared lock to read X
- If T2 wants to update X, an exclusive lock is required by T2
  - The exclusive lock is granted if and only if no other locks are held on the data item
  - If T1 has a shared lock on X. T2 can not get an exclusive lock on X – **mutual exclusive rule**
  - T2 must wait until T1 commits

41

## Lock Types

- A shared/exclusive lock schema increases the lock manager's overhead
  - The type of lock held must be known before a lock can be granted
  - Three lock operations exist:
    - READ\_LOCK to check the type of lock
    - WRITE\_LOCK to issue the lock
    - UNLOCK to release the lock
  - The schema has been enhanced to allow a lock upgrade from shared to exclusive and a lock downgrade from exclusive to shared

42

## Lock Types

- Although locks prevent serious data inconsistencies, they can lead to major problems
  - The resulting transaction schedule might not be serializable
  - The schedule might create deadlocks
    - A deadlock occurs when two transactions wait indefinitely for each other to unlock data
    - A database deadlock (similar to traffic gridlock) is caused when two or more transactions wait for each to unlock data

43

## Two-Phase Locking to Ensure Serializability

- Defines how transactions acquire and relinquish locks
- Guarantees serializability, but does not prevent deadlocks
  - Growing phase
    - Transaction acquires all required locks without unlocking any data
  - Shrinking phase
    - Transaction releases all locks and cannot obtain any new lock

44


## Two-Phase Locking to Ensure Serializability

- Governed by the following rules:
  - Two transactions cannot have conflicting locks
  - No unlock operation can precede a lock operation in the same transaction
  - No data are affected until all locks are obtained, meaning until the transaction is in its locked point

## Two-Phase Locking to Ensure Serializability

- The transaction first acquires the two locks it needs
- When it has two locks, it reaches its locked point
- Next, the data are modified to conform to the transaction requirements
- Finally, the transaction is completed and it releases all of the locks it acquired in the first phase
- Two phase locking increases transaction processing and cost and might cause deadlocks

FIGURE 10.7  
Two-phase locking protocol



SOURCE: Course Technology/Cengage Learning

## Deadlocks

- Condition that occurs when two transactions wait for each other to unlock data
  - T1 = access data items X and Y
  - T2 = access data items Y and X
    - If T1 has not unlocked data item Y, T2 cannot begin
    - If T2 has not unlocked data item X, T1 cannot continue
    - T1 and T2 each wait for the other to unlock the required data item – **deadly embrace**
- Possible only if one of the transactions wants to obtain an exclusive lock on a data item
  - No deadlock condition can exist among shared locks

TABLE 10.13 How a Deadlock Condition Is Created

| TIME | TRANSACTION | REPLY | LOCK STATUS |          |
|------|-------------|-------|-------------|----------|
| 0    |             |       | Data X      | Data Y   |
| 1    | T1:LOCK(X)  | OK    | Unlocked    | Unlocked |
| 2    | T2: LOCK(Y) | OK    | Locked      | Unlocked |
| 3    | T1:LOCK(Y)  | WAIT  | Locked      | Locked   |
| 4    | T2:LOCK(X)  | WAIT  | Locked      | Locked   |
| 5    | T1:LOCK(Y)  | WAIT  | Locked      | Locked   |
| 6    | T2:LOCK(X)  | WAIT  | Locked      | Locked   |
| 7    | T1:LOCK(Y)  | WAIT  | Locked      | Locked   |
| 8    | T2:LOCK(X)  | WAIT  | Locked      | Locked   |
| 9    | T1:LOCK(Y)  | WAIT  | Locked      | Locked   |
| ...  | .....       | ..... | .....       | .....    |
| ...  | .....       | ..... | .....       | .....    |
| ...  | .....       | ..... | .....       | .....    |
| ...  | .....       | ..... | .....       | .....    |



## Prevention

- A transaction requesting a new lock is aborted when there is the possibility that a deadlock can occur.
- If the transaction is aborted, all changes made by this transaction are rolled back and all locks obtained by the transaction are released
- The transaction is then rescheduled for execution
- It avoids the condition that leads to deadlocking

# Controlling Deadlock

## Detection

- The DBMS periodically tests the database for deadlocks
- If found, the “victim” transaction is aborted (rolled back and restarted) and the other transaction continues

## Avoidance

- The transaction must obtain all of the locks it needs before it can be executed
- This avoids the rolling back of conflicting transactions by requiring that locks be obtained in succession
- Serial lock assignment increases action response times

# Controlling Deadlock

# Deadlocks

- Choice of deadlock control method depends on database environment
  - Low probability of deadlock; detection recommended
  - High probability; prevention recommended
  - If response time is not high on the system’s priority list, deadlock avoidance might be employed
- All current DBMSs support deadlock detection while some use a blend of prevention and avoidance techniques for other types of data, such as data warehouses or XML data



## Concurrency Control with Time Stamping Methods

- Assigns global unique time stamp to each transaction
- Produces explicit order in which transactions are submitted to DBMS
- Properties of timestamps
  - Uniqueness
    - Ensures that no equal time stamp values can exist
  - Monotonicity
    - Ensures that time stamp values always increase

53

## Concurrency Control with Time Stamping Methods

- All database operations within the same transaction must have the same timestamp
- The DBMS executes conflicting operations in timestamp order, thereby ensuring serializability of the transactions
- If two transactions conflict, one is stopped, rolled back, rescheduled and assigned a new timestamp value

54

## Concurrency Control with Time Stamping Methods

- Disadvantage - each value stored in the database requires two additional timestamp fields; one for the last time the field was read and one for the last update
  - This increases memory needs and the database's processing overhead
- Timestamping demands a lot of system resources because many transactions might have to be stopped, rescheduled and restamped

55

## Wait/Die and Wound/Wait Schemes

- How to decide which transaction to rollback and which continues executing
  - **Wait/die**
    - Older transaction requests lock - it waits until younger is completes and locks released
    - Younger transaction requests lock - it dies (rolls back) and is rescheduled using the same timestamp
  - **Wound/wait**
    - Older transaction requests lock, it preempts (wounds) the younger transaction by rolling it back and rescheduling it using the same timestamp
    - Younger transaction requests lock, it will wait until the other transaction is completed and the locks released

56

- Wait for lock request can cause some transactions to wait indefinitely, causing a deadlock
- To prevent a deadlock, each lock request has an associated time-out value
  - If the lock is not granted before the time-out expires, the transaction is rolled back

**TABLE 10.14** Wait/Die and Wound/Wait Concurrency Control Schemes

| TRANSACTION REQUESTING LOCK | TRANSACTION OWNING LOCK | WAIT/DIE SCHEME  | WOUND/WAIT SCHEME   |
|-----------------------------|-------------------------|--|---|
| T1 (11548789)               | T2 (19562545)           | <ul style="list-style-type: none"> <li>• T1 waits until T2 is completed and T2 releases its locks.</li> </ul>                    | <ul style="list-style-type: none"> <li>• T1 preempts (rolls back) T2.</li> <li>• T2 is rescheduled using the same timestamp.</li> </ul> |
| T2 (19562545)               | T1 (11548789)           | <ul style="list-style-type: none"> <li>• T2 dies (rolls back).</li> <li>• T2 is rescheduled using the same timestamp.</li> </ul> | <ul style="list-style-type: none"> <li>• T2 waits until T1 is completed and T1 releases its locks.</li> </ul>                           |

57

Database Systems, 10th Edition

58

16/10/2022

16/10/2022

## Concurrency Control with Optimistic Methods

- Read Phase
  - Transaction reads the database, executes computations, makes updates to a private copy of the database values (temporary update file) not accessed by the remaining transactions
- Validation Phase
  - Transaction is validated to ensure that the changes made will not affect the integrity and consistency of the database
    - If validation test is positive, the transaction goes to the write phase
    - If validation test is negative, transaction is restarted and the changes discarded
- Write Phase
  - The changes are permanently applied to the database

59

16/10/2022

## Concurrency Control with Optimistic Methods

- Optimistic approach
  - Based on assumption that majority of database operations do not conflict
  - Does not require locking or time stamping techniques
  - Transaction is executed without restrictions until it is committed
  - Each transaction moves through two or three phases - read, validation, and write

## Concurrency Control with Optimistic Methods

- This approach is acceptable for most read or query database systems that require few update transactions
- In a heavily used DBMS environment, one or more of the above techniques will be used
- However, it may be necessary at times to employ database recovery techniques to restore the database to a consistent state

60

16/10/2022

# Database Recovery Management

- Restores database to previous consistent state
- Based on atomic transaction property
  - All portions of transaction are treated as single logical unit of work
  - All operations are applied and completed to produce consistent database
- If transaction operation cannot be completed:
  - Transaction aborted
  - Changes to database are rolled back

# Database Recovery Management

- In addition to recovering transactions, recovery techniques also apply to the database and system after some type of critical error
  - Hardware/software failures
  - Human caused incidents
  - Natural disasters

## Transaction Recovery

- Use data in the transaction log to recover a database from an inconsistent state to a consistent state
- Four concepts that affect the recovery process
  - **Write-ahead-log protocol:** ensures transaction logs are written before database data is updated
    - In case of a failure, the DB can be recovered using data in the transaction log
  - **Redundant transaction logs:** ensure physical disk failure will not impair ability to recover

## Transaction Recovery

- **Buffers:** temporary storage areas in primary memory
  - To improve processing time, the DBMS software reads the data from the physical disk and stores a copy of it on a “buffer” in primary memory
  - When a transaction updates data, it actually updates the copy of the data in the buffer because that’s faster than updating the physical disk
  - Later, all buffers that contain data are written to a physical disk during a single operation thereby saving significant processing time

- **Checkpoints:** operations in which DBMS writes all its updated buffers to disk
  - While this is happening, the DBMS does not execute any other requests
  - Scheduled by the DBMS several times per hour
  - A checkpoint operation is also registered in the transaction log
  - The physical database and the transaction log are in synch
    - Synching is required because update operations update the copy of the data in the buffers and not in the physical database

65

- **Deferred-write (deferred update) technique**
  - Only transaction log is updated
  - DB is physically updated only after the transaction reaches its commit point
  - If the transaction aborts before a commit, no rollback is needed

66

- Recovery process for all started and committed transactions (before the failure) follows these steps
  - identify last checkpoint (i.e., when transaction data was physically saved to disk)
  - If transaction committed before checkpoint:
    - Do nothing, the data was already saved
  - If transaction committed after checkpoint:
    - Use transaction log to redo the transaction using the “after” values
      - Changes are made in oldest to newest order
  - If transaction had ROLLBACK operation after the last checkpoint or that was left active (with neither a COMMIT nor a ROLLBACK) before the failure:
    - Do nothing, DB was never updated

67

- **Write-through technique**
  - Database is immediately updated by transaction operations during transaction’s execution even before the transaction reaches its commit point
- Recovery process: identify last checkpoint (when data were physically written to the disk)
  - If transaction committed before checkpoint:
    - Do nothing; data already saved
  - If transaction committed after last checkpoint:
    - DBMS redoes the transaction using “after” values in oldest to newest order
  - If transaction had ROLLBACK or was left active (with neither a COMMIT nor a ROLLBACK) before failure occurred:
    - DBMS uses the transaction log to ROLLBACK using the “before” values in newest to oldest order

68

TABLE  
10.15

A Transaction Log for Transaction Recovery Examples

| TRL ID | TRX NUM | PREV PTR | NEXT PTR | OPERATION  | TABLE                   | ROW ID   | ATTRIBUTE    | BEFORE VALUE | AFTER VALUE             |
|--------|---------|----------|----------|------------|-------------------------|----------|--------------|--------------|-------------------------|
| 341    | 101     | Null     | 352      | START      | ****Start Transaction   |          |              |              |                         |
| 352    | 101     | 341      | 363      | UPDATE     | PRODUCT                 | 54778-2T | PROD_QOH     | 45           | 43                      |
| 363    | 101     | 352      | 365      | UPDATE     | CUSTOMER                | 10011    | CUST_BALANCE | 615.73       | 675.62                  |
| 365    | 101     | 363      | Null     | COMMIT     | **** End of Transaction |          |              |              |                         |
| 397    | 106     | Null     | 405      | START      | ****Start Transaction   |          |              |              |                         |
| 405    | 106     | 397      | 415      | INSERT     | INVOICE                 | 1009     |              |              | 1009,10016, ...         |
| 415    | 106     | 405      | 419      | INSERT     | LINE                    | 1009,1   |              |              | 1009,1, 89-WRE-Q,1, ... |
| 419    | 106     | 415      | 427      | UPDATE     | PRODUCT                 | 89-WRE-Q | PROD_QOH     | 12           | 11                      |
| 423    |         |          |          | CHECKPOINT |                         |          |              |              |                         |
| 427    | 106     | 419      | 431      | UPDATE     | CUSTOMER                | 10016    | CUST_BALANCE | 0.00         | 277.55                  |
| 431    | 106     | 427      | 457      | INSERT     | ACCT_TRANSACTION        | 10007    |              |              | 1007,18-JAN-2012, ...   |
| 457    | 106     | 431      | Null     | COMMIT     | **** End of Transaction |          |              |              |                         |
| 521    | 155     | Null     | 525      | START      | ****Start Transaction   |          |              |              |                         |
| 525    | 155     | 521      | 528      | UPDATE     | PRODUCT                 | 2232/QWE | PROD_QOH     | 6            | 26                      |
| 528    | 155     | 525      | Null     | COMMIT     | **** End of Transaction |          |              |              |                         |

\*\*\*\*\* C \*R\*A\* S\* H \*\*\*\*\*

- Transaction 101 consists of two UPDATE statements that reduce the QOH for 54778-2T and increase the customer balance for customer 10011 for a credit sale of 2 units of product 54778-27
- Transaction 106 represents the credit sale of 1 unit of 89-WRE-Q to customer 10016 for \$277.55. This transaction consists of 5 DML statement – 3 INSERTs and 2 UPDATES
- Transaction 155 consists of 1 UPDATE that increases the QOH of 2232/QWE from 6 to 26
- A database checkpoint writes all updated database buffers for previously uncommitted transaction to disk. In this case, all changes made by 101.

# Transaction Recovery

- Using the deferred update method, the following will happen during a recovery
  - Identify the last checkpoint – TRL ID 423
  - Transaction 101 started and ended before the last checkpoint so no action needs to be taken
  - For each transaction committed after TRL ID 423, the DBMS uses the transaction log to write changes to disk using the “after” values. For transaction 106
    - Find COMMIT (TRL ID 457)
    - Use the previous pointer values to locate the start of the transaction (TRL ID 397)
    - Use the next pointer values to locate each DML and apply the “after” values (TRL ID 405, 415, 419, 427 and 431)
    - Repeat the process for transaction 155
  - Any other transaction will be ignored as no changes were written to disk