Reinforcement learning

Dr. Aissa Boulmerka

The National School of Artificial Intelligence aissa.boulmerka@ensia.edu.dz

2024-2025

Chapter 8 On-policy Prediction with Approximation

Outline

- Value-function Approximation
- Generalization and Discrimination
- The Prediction Objective (VE)
- Stochastic-gradient and Semi-gradient Methods
- Linear Methods
- Feature Construction for Linear Methods
- Nonlinear Function Approximation

Introduction

Introduction

- In this chapter, the approximate value function is represented not as a table but as a parameterized functional form with weight vector $\mathbf{w} \in \mathbb{R}^d$.
- We will write $\hat{v}(s, \mathbf{w}) \approx v_{\pi}(s)$ for the approximate value of state s given the weight vector \mathbf{w} . Note that \mathbf{w} is the vector of weights.
- For example, \hat{v} might be a linear function in features of the state. More generally, \hat{v} might be a non-linear function computed by a multi-layer artificial neural network.
- Typically, the number of weights (the dimensionality of **w**) is much less than the number of **states** ($d \ll |\mathcal{S}|$), and changing one weight changes the estimated value of many states.
- Consequently, when a single state is updated, the change generalizes from that state to affect the values of many other states. Such generalization makes the learning potentially more powerful but also potentially more difficult to manage and understand.

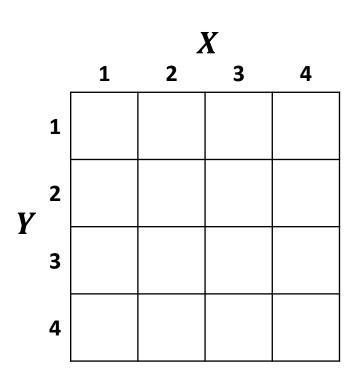
Value-Function Approximation

Parameterized value-function

$$\hat{v}(s, \mathbf{w}) \approx v_{\pi}(s)$$

Example of parameterized value-function

$$\hat{v}(s, \mathbf{w}) \doteq w_1 X + w_2 Y$$
We only have to store two weights w_1 and w_2



$$w_1 = 1$$
 $w_2 = 1$

		\boldsymbol{X}		
	1	2	3	4
1	2	3	4	5
2 Y	3	4	5	6
3	4	5	6	7
4	5	6	7	8

w_1	=	2	
w_2	=	1	

		\boldsymbol{X}			
	ı	1	2	3	4
Y	1	3	5	7	9
	2	4	6	8	10
	3	5	7	9	11
	4	6	8	10	12

$$w_1 = -1$$

$$w_2 = 1$$

		\boldsymbol{X}			
	ı	1	2	3	4
Y	1	0	-1	-2	-3
	2	1	0	-1	-2
	3	2	1	0	-1
	4	3	2	1	0

$$w_1 = 4$$
 $w_2 = 1$

		\boldsymbol{X}			
	ı	1	2	3	4
Y	1	5	9	13	17
	2	6	10	14	18
1	3	7	11	15	19
	4	8	12	16	20

Linear value-function approximation

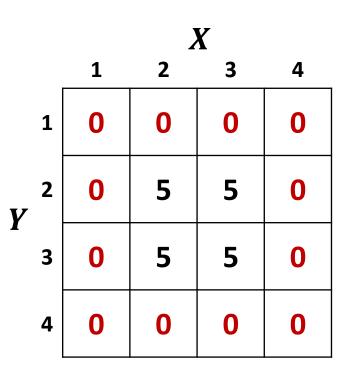
$$\hat{v}(s, \mathbf{w}) \doteq \sum w_i x_i(s) = \langle \mathbf{w}, \mathbf{x}(s) \rangle$$

- The value of each state is represented by a linear function of the weights.
- This simply means that the value of each state, is computed as the sum of the weights multiplied by some fixed attributes of the state called features.
- < w, x(s) > is the **dot product (inner product)** of **weight vector w** and the **feature vector** x(s).

Limitations of Linear Value Function Approximation

$$\hat{v}(s, \mathbf{w}) \doteq \sum w_i x_i(s)$$
?

- We cannot represent this as a linear function of X and Y.
- Here X and Y are not good features for this problem.
- There are many powerful methods to construct features.



Tabular value-functions are linear functions

State	Value
S_1	
S_2	
S_3	
•••	
S_i	
•••	
S ₁₆	

- Linear function approximation is actually very general.
- In fact, even a tabular representation is a special case of linear function approximation.
- How to implement that?

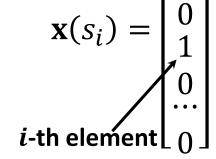
Tabular value-functions are linear functions

State	Value
S_1	w_1
S_2	w_2
S_3	w_3
•••	•••
S_i	w_i
•••	•••
S ₁₆	w ₁₆

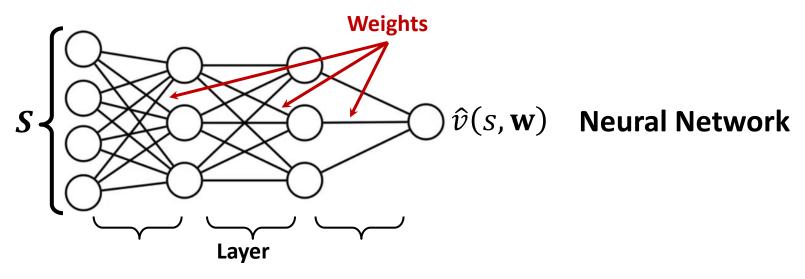
- Let's choose our features to be indicator functions for particular states.
- For state s_i , feature i is one and the remaining features are 0.
- We have 16 features, one for each state.

$$\hat{v}(s, \mathbf{w}) \doteq \langle \mathbf{w}, \mathbf{x}(s) \rangle$$

= w_i



Nonlinear function approximation



- Neural networks are an example of a nonlinear function of state. The output of the network is an approximate value for a given state.
- The state is passed to the network as the input. All the connections in the network correspond to real valued weights.
- This process transforms the input state through a sequence of layers to finally produce the value estimate.

Generalization and Discrimination

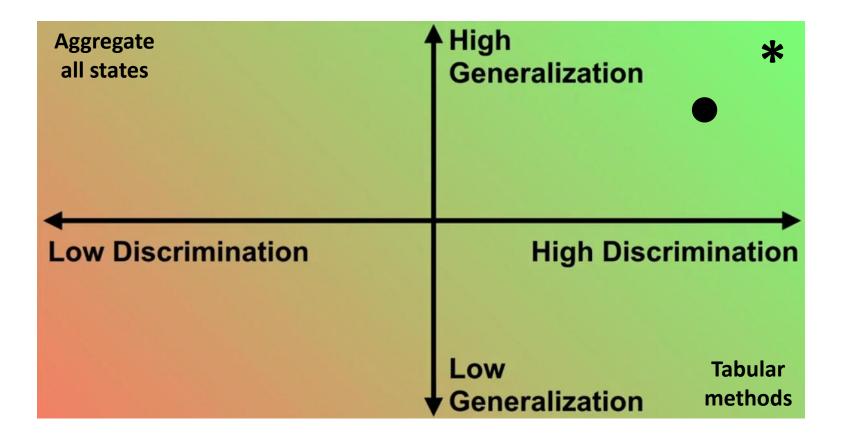
Generalization

- Generalization intuitively means applying knowledge about specific situations to draw conclusions about a wider variety of situations.
- Generalization, in the context of policy evaluation, means that updates to the value estimate of one state influence the value of other states.
- Imagine a robot, observing the world through a set of distance sensors. In many locations, it would take the same amount of time to drive to the nearest object. Even though they correspond to different sensor readings, these locations have similar values. Thus, we might want the value function to generalize across those states.
- Generalization can speed learning by making better use of the experience we have.
- You may not have to visit every state as much to get this values correct if we can learn its value from similar states.

Discrimination

- On the other hand, discrimination means the ability to make the values for two states different to distinguish between the values for these two states.
- Going back to the example of a robot, imagine it is in a state where an object is three feet away, but behind a wall.
- Compare this to a state where an object is three feet away, but with a clear paths to reach it.
- The robot would want to assign different values to these states.
- So while it is useful to generalize between states with similar distance to the nearest object, it is also important that we discriminate between states based on other information when it is likely to impact their value.

Generalization and Discrimination



Frame value estimation as supervised learning

- Supervised learning methods learn a function from an offline dataset. This is very different from reinforcement learning. But supervised learning methods can be useful for handling parts of the reinforcement learning problem.
- In reinforcement learning, an agent interacts with an environment and continually generates new data. This is often called the online setting. The proposed function approximation technique should work in the online setting.
- TD methods introduce an additional complication when applying techniques from supervised learning. TD methods use bootstrapping, meaning that our targets now depend on our own estimates.
- These estimates change as learning progresses. So our targets continually change. This is different than supervised learning where we have access to a ground truth label as the target.

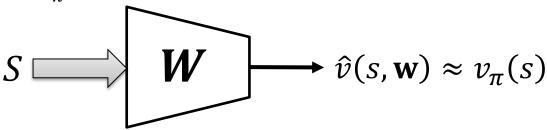
The Prediction Objective

The Prediction Objective \overline{VE}

 Suppose the following idealized scenario. We get a sequence of pairs of states and true values.

$$\{(S_1, v_{\pi}(S_1)), (S_2, v_{\pi}(S_2)), (S_3, v_{\pi}(S_3)), \cdots\}$$

• We want to use this **data** to find a **parameterized function** that closely **approximates** v_{π} .



- We will do this by adjusting the weights so that the output of the function approximate the associated value for a given state.
- To make our goal precise, we need to specify some measure of how close our approximation is to the value function.

The Mean Squared Value Error Objective

- Let's specify a state distribution $\mu(s) \ge 0$, $\sum_s \mu(s) = 1$, representing how much we care about the error in each state s.
- The error in a state s is computed using the square of the difference between the approximate value $\hat{v}(s, \mathbf{w})$ and the true value $v_{\pi}(s)$.
- Weighting this over the state space by μ , we obtain a natural **objective function**, the **Mean Squared Value Error**, denoted \overline{VE} :

$$\overline{VE} \doteq \sum_{s \in S} \mu(s) [v_{\pi}(s) - \hat{v}(s, \mathbf{w})]^2$$
.

- Often $\mu(s)$ is chosen to be the **fraction of time spent** in s. Under on-policy training this is called the **on-policy distribution**.
- The goal of defining this objective is to adapt the weights to minimize the mean squared value error.

Stochastic-Gradient and Semi-Gradient Methods

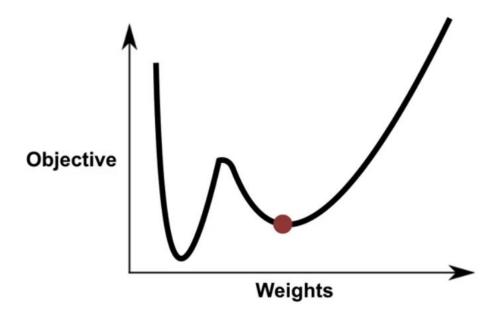
Stochastic-gradient methods

- Consider the weight vector $w = (w_1, w_2, \dots, w_d)^T$, and the approximate value function $\hat{v}(s, \mathbf{w})$ is a differentiable function of \mathbf{w} for all $s \in \mathcal{S}$.
- We will update **w** at each of a series of **discrete time steps**, $t = 0, 1, 2, 3, \cdots$. Also, consider w_t for the weight vector at each step t.
- Stochastic gradient-descent (SGD) methods are particularly well suited to online reinforcement learning. Stochastic gradient-descent (SGD) methods minimize error on the observed examples by adjusting the weight vector after each example by a small amount in the direction that most reduce the error on that example:

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t - \frac{1}{2} \alpha \nabla [v_{\pi}(S_t) - \hat{v}(S_t, \mathbf{w}_t)]^2$$
$$= \mathbf{w}_t + \alpha [v_{\pi}(S_t) - \hat{v}(S_t, \mathbf{w}_t)] \nabla \hat{v}(S_t, \mathbf{w}_t)$$

From gradient descent to stochastic gradient descent

$$\sum_{s \in \mathcal{S}} \mu(s) [v_{\pi}(s) - \hat{v}(s, \mathbf{w})] \nabla \hat{v}(s, \mathbf{w})$$
$$(S_{1}, v_{\pi}(S_{1})), (S_{2}, v_{\pi}(S_{2})), (S_{3}, v_{\pi}(S_{3})), \cdots$$



Gradient of the MSVE objective

$$\nabla \sum_{s \in S} \mu(s) [v_{\pi}(s) - \hat{v}(s, w)]^2$$

$$= \sum_{s \in S} \mu(s) \nabla [v_{\pi}(s) - \hat{v}(s, \mathbf{w})]^2$$

$$= -\sum_{s \in \mathcal{S}} \mu(s) 2[v_{\pi}(s) - \hat{v}(s, \mathbf{w})] \nabla \hat{v}(s, \mathbf{w})$$

Linear value function

$$\hat{v}(s, \mathbf{w}) \doteq \langle \mathbf{w}, \mathbf{x}(s) \rangle$$

$$\nabla \hat{v}(s, \mathbf{w}) = \mathbf{x}(s)$$

$$\nabla \mathbf{w} \propto \sum_{s \in S} \mu(s) [v_{\pi}(s) - \hat{v}(s, \mathbf{w})] \nabla \hat{v}(s, \mathbf{w})$$

Gradient Monte Carlo

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha [\mathbf{v}_{\pi}(S_t) - \hat{v}(S_t, \mathbf{w}_t)] \nabla \hat{v}(S_t, \mathbf{w}_t)$$

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha [G_t - \hat{v}(S_t, \mathbf{w}_t)] \nabla \hat{v}(S_t, \mathbf{w}_t)$$

Recall that :

$$v_{\pi}(s) \doteq \mathbb{E}_{\pi}[G_t | S_t = s]$$

The expectation of the gradient when we use a sampled return in place of the true value, is still equal to the gradient of the MSVE.

$$\mathbb{E}_{\pi} \left[\mathbf{v}_{\pi}(S_t) - \hat{v}(S_t, \mathbf{w}_t) \right] \nabla \hat{v}(S_t, \mathbf{w}_t)$$

$$= \mathbb{E}_{\pi} \left[\mathbf{G}_t - \hat{v}(S_t, \mathbf{w}_t) \right] \nabla \hat{v}(S_t, \mathbf{w}_t)$$

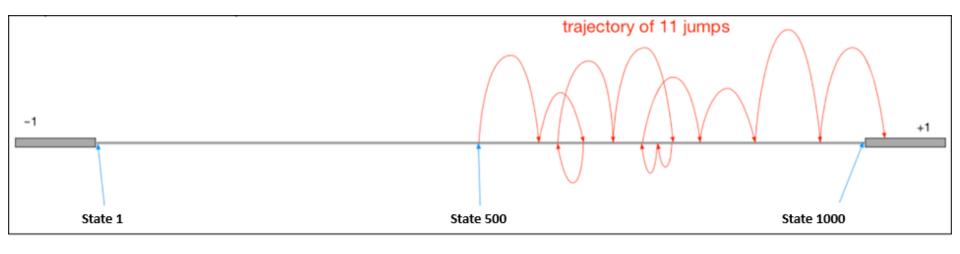
Gradient Monte Carlo state-value prediction

- Suppose the states in the examples are the states generated by **interaction** with the **environment** using **policy** π .
- The SGD method converges to a **locally optimal approximation** to $v_{\pi}(S_t)$.
- So, the gradient-descent version of Monte Carlo state-value prediction is guaranteed to find a locally optimal solution.

Gradient Monte Carlo Algorithm for Estimating $\hat{v} \approx v_{\pi}$

```
Input: the policy \pi to be evaluated
Input: a differentiable function \hat{v}: \mathbb{S} \times \mathbb{R}^d \to \mathbb{R}
Algorithm parameter: step size \alpha > 0
Initialize value-function weights \mathbf{w} \in \mathbb{R}^d arbitrarily (e.g., \mathbf{w} = \mathbf{0})
Loop forever (for each episode):
Generate an episode S_0, A_0, R_1, S_1, A_1, \dots, R_T, S_T using \pi
Loop for each step of episode, t = 0, 1, \dots, T - 1:
\mathbf{w} \leftarrow \mathbf{w} + \alpha \left[ G_t - \hat{v}(S_t, \mathbf{w}) \right] \nabla \hat{v}(S_t, \mathbf{w})
```

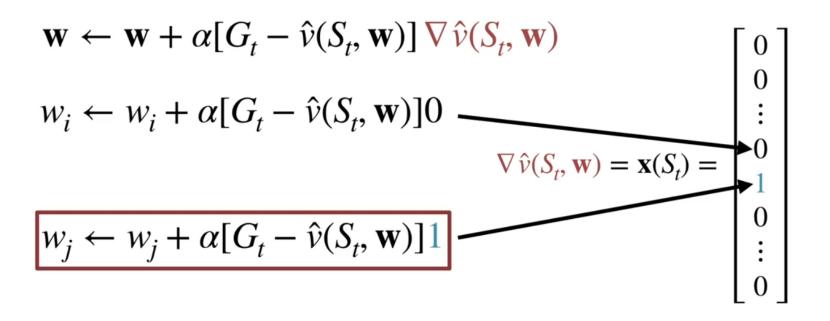
- Consider a 1000-state version of the random walk task. The states are numbered from 1 to 1000, left to right, and all episodes begin near the center, in state 500.
- State transitions are from the current state to one of the 100 neighboring states to its left, or to one of the 100 neighboring states to its right, all with equal probability.
- If the current state is **near an edge**, then there may be fewer than 100 neighbors on that side of it. In this case, all the probability that would have gone into those missing neighbors goes into the **probability of terminating** on that side (example, state 1 has a 0.5 chance of terminating on the left, and state 950 has a 0.25 chance of terminating on the right).
- Termination on the **left** produces a **reward** of -1, and termination on the **right** produces a **reward** of +1. All other **transitions** have a **reward** of zero.



State aggregation

State	Value		
s_1	3	[]	
s_2	3	$\left \left(\mathbf{v}(\mathbf{c}) - 1\right \right $	$\hat{v}(s, \mathbf{w}) = w_1$
s_3	3	$\left \begin{cases} \mathbf{x}(s) = \begin{bmatrix} 1 \\ 0 \end{bmatrix} \right $	$v(s, \mathbf{w}) - w_1$
s_4	3	۱) ا	
s_5	0) _[0]	
s_6	0	$\left \begin{cases} \mathbf{x}(s) = \begin{bmatrix} 0 \\ 1 \end{bmatrix} \right $	$\hat{v}(s, \mathbf{w}) = w_2$
<i>S</i> ₇	0	$\begin{bmatrix} \mathbf{A}(3) - 1 \end{bmatrix}$	$v(s, \mathbf{w}) - w_2$
s_8	0) r ₋₁	

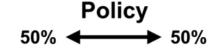
How to compute the gradient for Monte Carlo with state aggregation?



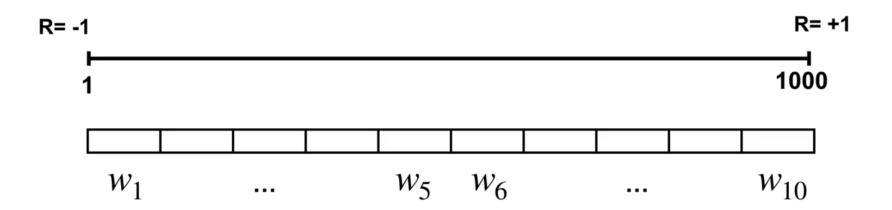
Monte Carlo update for a single episode

Return: 1, 1, 1, ..., 1

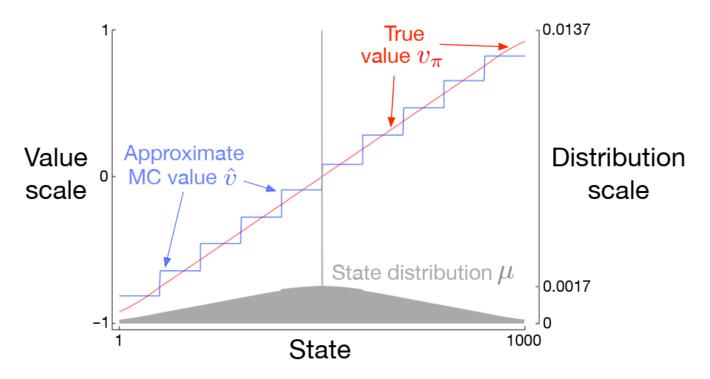
Visited states: 500, 423, 482, ..., 936



$$\alpha = 2 * 10^{-5}$$



Example: Random Walk



Function approximation by state aggregation on the 1000-state random walk task, using the gradient Monte Carlo algorithm.

The TD update for function approximation

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [U_t - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

• We can also replace U_t with a **bootstrap target**, such as the **one step TD target**.

$$U_t \doteq R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w})$$

- The TD target uses our current value estimate, which will likely not equal the true value function.
- Because of this we cannot guarantee this algorithm will converge to a local minimum of the value error.

$$\begin{split} \nabla \frac{1}{2} [U_t - \hat{v}(S_t, \mathbf{w})]^2 \\ \text{We have } U_t &\doteq R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) \\ &= (U_t - \hat{v}(S_t, \mathbf{w})) \big(\nabla U_t - \nabla \hat{v}(S_t, \mathbf{w}) \big) \\ \neq [U_t - \hat{v}(S_t, \mathbf{w}_t)] \nabla \hat{v}(S_t, \mathbf{w}_t) \text{ (The TD update)} \\ \text{This is true only if } \nabla U_t &= 0 \end{split}$$

For TD:

$$\nabla U_t = \nabla (R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}))$$
$$= \gamma \nabla \hat{v}(S_{t+1}, \mathbf{w})$$
$$\neq \mathbf{0}$$

- Bootstrapping methods are not in fact instances of **true gradient descent**. They take into account the effect of **changing the weight vector** \mathbf{w}_t on the estimate, but ignore its effect on the **target**. They include only a **part of the gradient** and, accordingly, we call them **semi-gradient methods**.
- Although semi-gradient (bootstrapping) methods do not converge as robustly as gradient methods, they do converge reliably in important cases such as the linear case.
- One advantage for this is that they typically enable significantly faster learning.
- Another advantage, is that they enable learning to be continual and online, without waiting for the end of an episode. This enables them to deal with continuing problems and provides computational advantages.

Semi-gradient TD(0) for estimating $\hat{v} \approx v_{\pi}$ Input: the policy π to be evaluated Input: a differentiable function $\hat{v}: \mathbb{S}^+ \times \mathbb{R}^d \to \mathbb{R}$ such that $\hat{v}(\text{terminal}, \cdot) = 0$ Algorithm parameter: step size $\alpha > 0$ Initialize value-function weights $\mathbf{w} \in \mathbb{R}^d$ arbitrarily (e.g., $\mathbf{w} = \mathbf{0}$) Loop for each episode: Initialize SLoop for each step of episode: Choose $A \sim \pi(\cdot|S)$ Take action A, observe R, S' $\mathbf{w} \leftarrow \mathbf{w} + \alpha [R + \gamma \hat{v}(S', \mathbf{w}) - \hat{v}(S, \mathbf{w})] \nabla \hat{v}(S, \mathbf{w})$

 $S \leftarrow S'$

until S is terminal

Comparing TD and MC with state aggregation

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [\mathbf{U}_t - \hat{v}(S_t, \mathbf{w})] \nabla \hat{v}(S_t, \mathbf{w})$$

Gradient Monte Carlo	Semi-Gradient TD
$Target: \pmb{U_t} = \pmb{G_t}$	$Target: \boldsymbol{U}_t = \boldsymbol{R}_{t+1} + \gamma \widehat{\boldsymbol{v}}(\boldsymbol{S}_{t+1}, \mathbf{w})$
 Gradient MC uses an unbiased estimate of the gradient of the value error. 	The update could be biased because the estimate in the target may not be accurate.
 It will approach a local minimum of the Mean Squared Value Error with more and more samples. 	 Since the value approximation will never be perfect even in the limit, the target may remain biased.
 Gradient Monte Carlo will converge to a local minimum of the mean squared value error. 	 Semi-gradient TD cannot guarantee to converge to a local minimum at the Mean Squared value error.

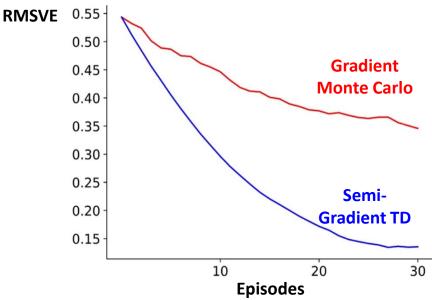
Comparing TD and MC with state aggregation

Experiment settings:

- 30 episodes
- 100 evenly spaced values of α between 0 and 1 with each algorithm.
- The best α for TD : 0.22 and for MC : 0.01.

Conclusions:

We can conclude that TD often learns
 faster than Monte Carlo. This is because TD
 can learn during the episode and has lower variance
 updates.



- Monte Carlo is better on long-run performance, it's not always the main concern.
- We can never run our experiments to achieve asymptotic performance.
- Early learning is perhaps more important in practice.

Linear Methods

Linear function approximation

- One of the most important special cases of **function approximation** is that in which the approximate function, $\hat{v}(s, \mathbf{w})$, is a **linear function** of the **weight** vector, \mathbf{w} .
- Corresponding to every state s, there is a real-valued vector

$$\mathbf{x}(s) \doteq (x_1(s), x_2(s), \cdots, x_d(s))^T$$

Linear methods approximate state-value function by the inner product between \mathbf{w} and $\mathbf{x}(s)$:

$$\hat{v}(s, \mathbf{w}) \doteq w^T \mathbf{x}(s) \doteq \sum_{i=1}^d w_i x_i(s)$$
.

In this case the approximate value function is said to be linear in the weights, or simply linear.

Linear function approximation

- The vector x(s) is called a feature vector representing state s.
- Each component $x_i(s)$ of $\mathbf{x}(s)$ is the value of a function $x_i: S \to \mathbb{R}$.
- We think of a feature as the entirety of one of these functions, and we call its value for a state s a feature of s.
- For linear methods, features are basis functions because they form a linear basis for the set of approximate functions.
- Constructing d-dimensional feature vectors to represent states is the same as selecting a set of d basis functions.
- Features may be defined in many different ways; we cover a few possibilities later in this course.

SGD update of linear function approximation

It is natural to use SGD updates with linear function approximation. The gradient of the approximate value function with respect to w in this case is:

$$\nabla \hat{v}(s, \mathbf{w}) = \mathbf{x}(s).$$

In this case the general SGD update reduces to a particularly simple form:

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha [U_t - \hat{v}(S_t, \mathbf{w}_t)] \mathbf{x}(S_t)$$

- Because it is so simple, the linear SGD case is one of the most favorable for mathematical analysis.
- Almost all useful convergence results for learning systems of all kinds are for linear (or simpler) function approximation methods.

TD update of linear function approximation

- The semi-gradient TD(0) algorithm also converges under linear function approximation. The weight vector converged to is also not the global optimum, but rather a point near the local optimum.
- The update at each time t is

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha \delta_t \mathbf{x}(S_t)$$
$$\delta_t = R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t)$$

Where δ_t is the **TD error**.

This fixed basis given by the expert design features has a large impact on the update. If well-designed, we can get effective value function approximation with a simple update.

Tabular TD is a special case of linear TD

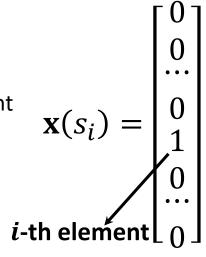
We can show that linear TD is a strict generalization of both tabular TD and TD with state aggregation.

$$\mathbf{w} \leftarrow \mathbf{w} + \alpha [R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})] \mathbf{x}(S_t)$$

• In the update, the **feature** vector $\mathbf{x}(S_t)$ selects a single weight associated with the current state.

$$\mathbf{w}_i \leftarrow \mathbf{w}_i + \alpha [R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}) - \hat{v}(S_t, \mathbf{w})]$$

 We can use the same analysis to show that TD with state aggregation is also a special case of linear TD.



$$\hat{v}(S_i, \mathbf{w}) = w_i$$

The Expected TD update

• We can expand the **TD update** like the following formula, note that we have used the notational shorthand $\mathbf{x}_t = \mathbf{x}(S_t)$.

$$\mathbf{w}_{t+1} \doteq \mathbf{w}_t + \alpha [R_{t+1} + \gamma \hat{v}(S_{t+1}, \mathbf{w}_t) - \hat{v}(S_t, \mathbf{w}_t)] \mathbf{x}_t$$

$$= \mathbf{w}_t + \alpha [R_{t+1} + \gamma \mathbf{w}_t^T \mathbf{x}_{t+1} - \mathbf{w}_t^T \mathbf{x}_t] \mathbf{x}_t$$

$$= \mathbf{w}_t + \alpha [R_{t+1} \mathbf{x}_t - \mathbf{x}_t (\mathbf{x}_t - \gamma \mathbf{x}_{t+1})^T \mathbf{w}_t]$$

■ The **expected update** characterizes the **expected change** in the weight from one time step to the next:

$$\mathbb{E}[\Delta \mathbf{w}_t] = \alpha(\mathbf{b} - \mathbf{A}\mathbf{w}_t)$$

Where
$$\mathbf{b} \doteq \mathbb{E}[R_{t+1}\mathbf{x}_t] \in \mathbb{R}^d$$
 and $\mathbf{A} \doteq \mathbb{E}[\mathbf{x}_t(\mathbf{x}_t - \gamma \mathbf{x}_{t+1})^T] \in \mathbb{R}^d \times \mathbb{R}^d$

The TD fixed point

It is clear that, if the system **converges**, it must converge to the weight vector \mathbf{w}_{TD} at which

$$\mathbb{E}[\Delta \mathbf{w}_{TD}] = \alpha(\mathbf{b} - \mathbf{A}\mathbf{w}_{TD}) = 0$$
$$\Rightarrow \mathbf{w}_{TD} = \mathbf{A}^{-1}\mathbf{b}$$

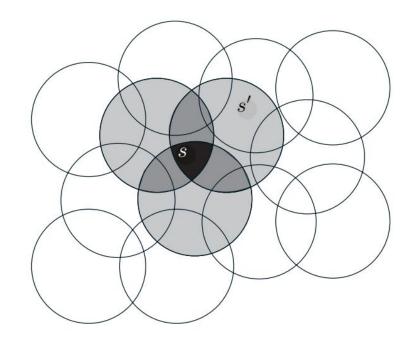
- If A is invertible, w_{TD} is a solution to this linear system. We call this solution the TD fixed point. In fact linear semi-gradient TD(0) converges to this point.
- \mathbf{w}_{TD} minimizes $(\mathbf{b} \mathbf{A}\mathbf{w})^T(\mathbf{b} \mathbf{A}\mathbf{w})$. This **objective** extends the connection between **TD** and Bellman equations, to the function approximation setting.
- At the TD **fixed point**, the \overline{VE} is **bounded**

$$\overline{VE}(\mathbf{w}_{TD}) \leq \frac{1}{1-\gamma} \min_{\mathbf{w}} \overline{VE}(\mathbf{w})$$

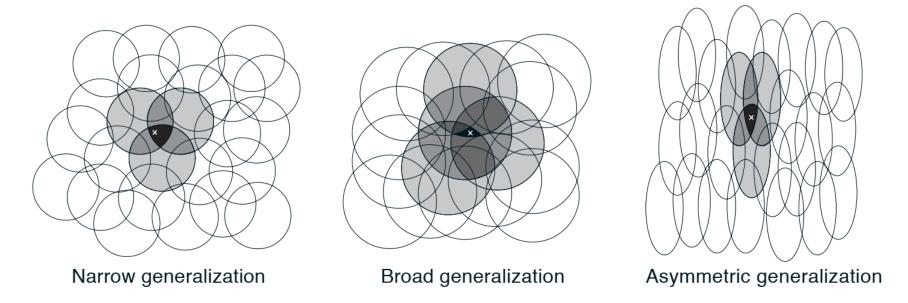
Feature Construction for Linear Methods

Coarse coding

- Features corresponding to circles in state space.
- If the state is inside a circle, then the corresponding feature has the value 1 and is said to be present; otherwise the feature is 0 and is said to be absent.
- This kind of 1–0-valued feature is called a binary feature.
- Representing a state with features that overlap in this way is known as coarse coding.

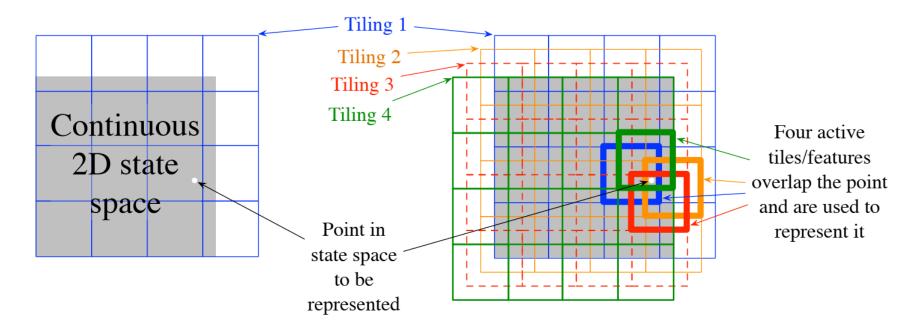


Coarse coding generalization



- Generalization in linear function approximation methods is determined by the sizes and shapes of the features' receptive fields.
- All three of these cases have roughly the same number and density of features.

Tile coding

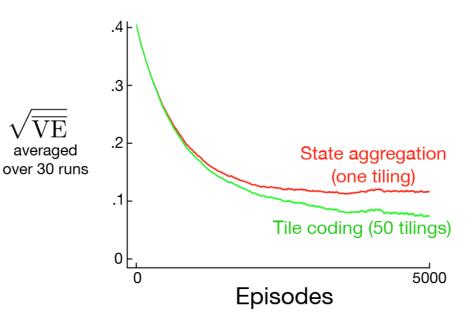


- In tile coding the receptive fields of the features are grouped into partitions of the state space.
- Each such partition is called a tiling, and each element of the partition is called a tile.

Why we use coarse coding?

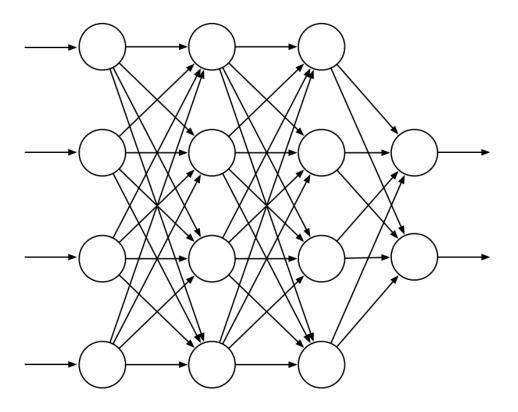
1000-state random walk example for the gradient Monte Carlo algorithm with a single tiling and with multiple tilings.

- The space of 1000 states was treated as a single continuous dimension, covered with tiles each 200 states wide.
- The multiple tilings were offset from each other by 4 states.
- The step-size parameter was set so that the initial learning rate in the two cases was the same,
 - $\alpha = 0.0001$ for the single tiling
 - $\alpha = 0.0001/50$ for the 50 tilings.

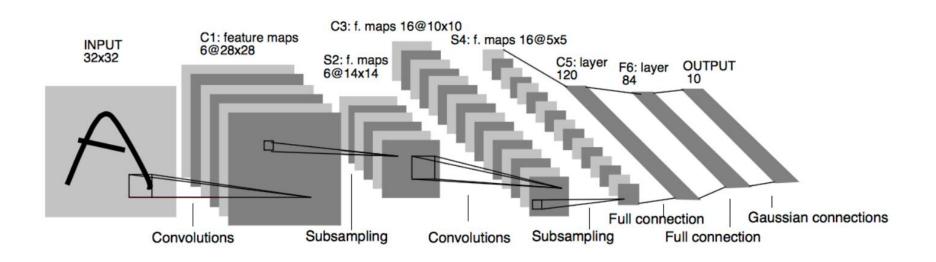


Nonlinear Function Approximation

Artificial Neural Network



Deep Neural Network



Deep Neural Network

- In theory, a neural network need not be deep. A neural network with a single hidden layer can approximate any continuous function given that is sufficiently wide.
- We call this the universal approximation property.
- Practical experience and theory suggests that deep neural networks may make it easier to approximate complex functions.
- One reason for this is that the depth allows composition of features.
- Composition can produce more specialized features by combining modular components.
- Overall, depth in a network can significantly improve our agent's ability to learn features.

Thank you! Q/A